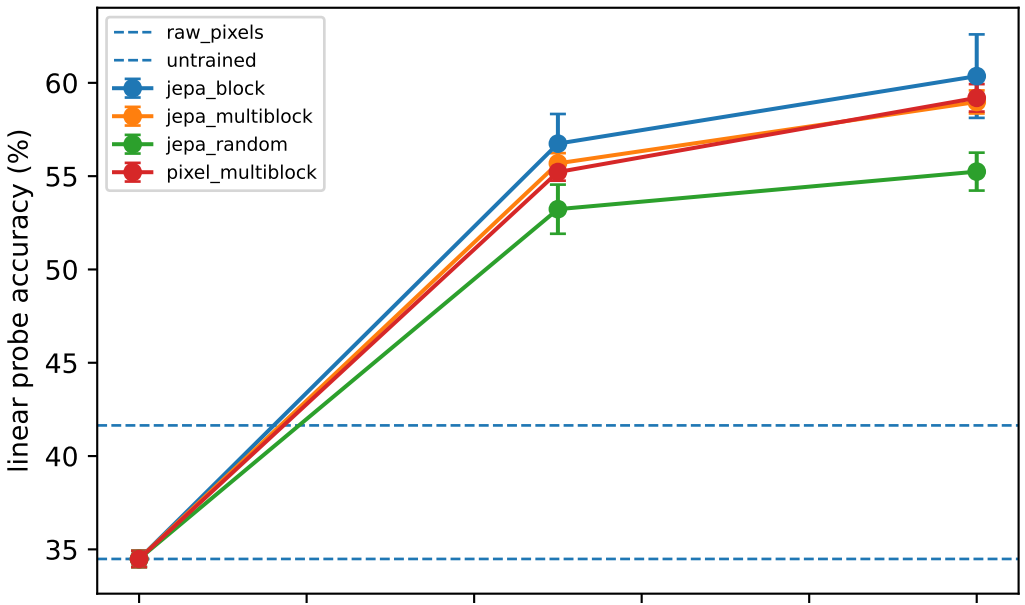
cifar10: accuracy vs training (mean +/- std over seeds)
raw_pixels
untrained
jepa_block
jepa_multiblock
jepa_random
pixel_multiblock
linear probe accuracy (%)
pretraining epoch
0
10
20
30
40
50
35
40
45
50
55
60

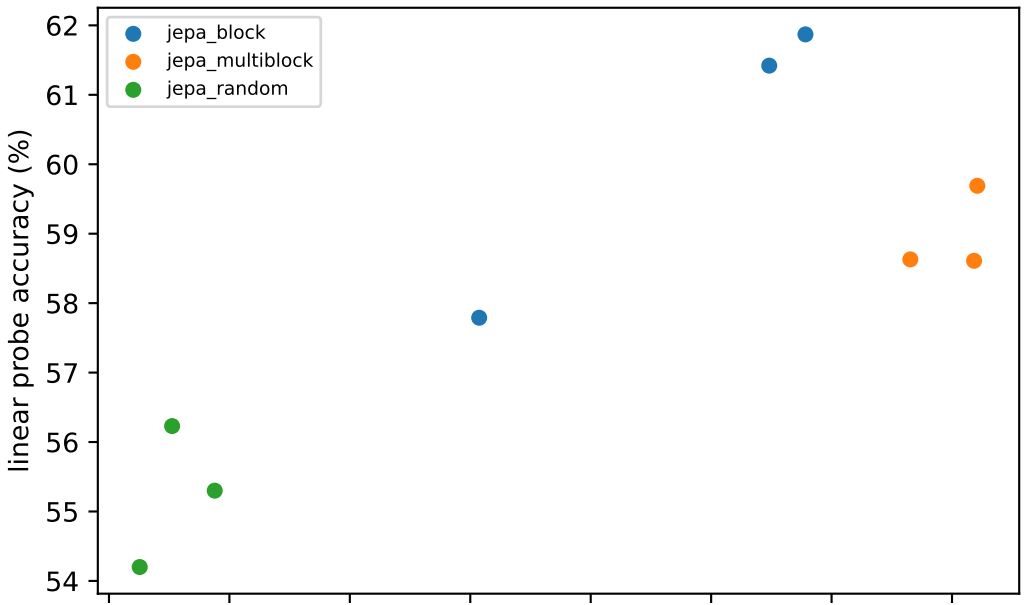
does training loss predict accuracy? (one dot per run)
jepa_block
jepa_multiblock
jepa_random
linear probe accuracy (%)
final training loss (latent-target runs)
0.075
0.100
0.125
0.150
0.175
0.200
0.225
0.250
54
55
56
57
58
59
60
61
62